

## Lab 4 - Part 2: Document Classification, Sentiment Analysis & Topic Modeling

**Course:** Natural Language Processing

**Objectives:**

- Build document classifiers (intro + advanced)
- Perform sentiment analysis on different domains
- Discover topics using unsupervised learning
- Compare different feature extraction methods

### Instructions

1. Complete all exercises marked with `# YOUR CODE HERE`
2. Answer all **written questions** in the designated markdown cells
3. Save your completed notebook
4. Push to your Git repository and send the link to: [yoroba93@gmail.com](mailto:yoroba93@gmail.com)

### Personal Analysis Required

This lab contains questions requiring YOUR personal interpretation.

### Use Cases Covered

Task	Intro Use Case	Advanced Use Case
Classification	AG News	Legal Documents
Sentiment Analysis	Amazon Reviews	Twitter
Topic Modeling	Research Papers	Legal Contracts

## Setup

```
1 # Install required libraries (uncomment if needed)
2 # !pip install datasets scikit-learn nltk pandas numpy matplotlib seaborn wordcloud gensim
```

```
1 import nltk
2
3 nltk.download('punkt')
4 nltk.download('punkt_tab')
5 nltk.download('stopwords')
6 nltk.download('wordnet')
7

[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Package punkt is already up-to-date!
[nltk_data] Downloading package punkt_tab to /root/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt_tab.zip.
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data]   Package wordnet is already up-to-date!
True
```

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from collections import Counter
6 import re
7 import warnings
8 warnings.filterwarnings('ignore')
9
10 # NLTK
11 import nltk
12 nltk.download('punkt', quiet=True)
13 nltk.download('stopwords', quiet=True)
14 nltk.download('wordnet', quiet=True)
15 from nltk.corpus import stopwords
16 from nltk.tokenize import word_tokenize
17 from nltk.stem import WordNetLemmatizer
18
```

```
19 # Scikit-learn
20 from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer
21 from sklearn.model_selection import train_test_split
22 from sklearn.naive_bayes import MultinomialNB
23 from sklearn.linear_model import LogisticRegression
24 from sklearn.svm import LinearSVC
25 from sklearn.ensemble import RandomForestClassifier
26 from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, f1_score
27 from sklearn.decomposition import LatentDirichletAllocation, NMF
28 from sklearn.pipeline import Pipeline
29
30 # Hugging Face datasets
31 from datasets import load_dataset
32
-- . . . . .
```

```
1 # Common preprocessing function
2 stop_words = set(stopwords.words('english'))
3 lemmatizer = WordNetLemmatizer()
4
5 def preprocess_simple(text):
6     """Basic preprocessing: lowercase, remove punctuation."""
7     text = str(text).lower()
8     text = re.sub(r'[^a-zA-Z\s]', '', text)
9     return ' '.join(text.split())
10
11 def preprocess_advanced(text):
12     """Advanced preprocessing: lowercase, remove punct, stopwords, lemmatize."""
13     text = str(text).lower()
14     text = re.sub(r'[^a-zA-Z\s]', '', text)
15     tokens = word_tokenize(text)
16     tokens = [lemmatizer.lemmatize(t) for t in tokens if t not in stop_words and len(t) > 2]
17     return ' '.join(tokens)
18
19 print("Preprocessing functions ready!")
```

Preprocessing functions ready!

## ▼ PART A: Document Classification

We will work with two use cases:

1. **Intro:** News Topic Classification (AG News)
  2. **Advanced:** Legal Document Classification (LexGLUE)

## ▼ A.1 Intro: News Topic Classification (AG News)

**Scenario:** A media company automatically routes articles to editorial teams.

## Feature Extraction: TF-IDF

```
1 # Load AG News dataset
2 print("Loading AG News dataset...")
3 ag_news = load_dataset("ag_news")
4
5 # Use subset for faster processing
6 ag_train = pd.DataFrame(ag_news['train']).sample(n=8000, random_state=42)
7 ag_test = pd.DataFrame(ag_news['test']).sample(n=2000, random_state=42)
8
9 # Label mapping
10 ag_labels = {0: 'World', 1: 'Sports', 2: 'Business', 3: 'Sci/Tech'}
11 ag_train['label_name'] = ag_train['label'].map(ag_labels)
12 ag_test['label_name'] = ag_test['label'].map(ag_labels)
13
14 print(f"Train: {len(ag_train)}, Test: {len(ag_test)}")
15 print(f"\nCategories: {list(ag_labels.values())}")
16 print(ag_train['label name'].value_counts())
```

```

Loading AG News dataset...
README.md: 8.07k? [00:00<00:00, 127kB/s]

data/train-00000-of-00001.parquet: 100% 18.6M/18.6M [00:01<00:00, 18.3MB/s]
data/test-00000-of-00001.parquet: 100% 1.23M/1.23M [00:00<00:00, 2.01MB/s]

Generating train split: 100% 120000/120000 [00:00<00:00, 213693.61 examples/s]
Generating test split: 100% 7600/7600 [00:00<00:00, 42237.01 examples/s]

Train: 8000, Test: 2000

Categories: ['World', 'Sports', 'Business', 'Sci/Tech']
label_name
Sports      2074
Sci/Tech    2021
Business    1959
World       1946
Name: count, dtype: int64

```

```

1 # Preprocess
2 ag_train['text_clean'] = ag_train['text'].apply(preprocess_simple)
3 ag_test['text_clean'] = ag_test['text'].apply(preprocess_simple)
4
5 # TF-IDF Vectorization
6 #from sklearn.feature_extraction.text import TfidfVectorizer
7
8 # TF-IDF Vectorization
9 tfidf_ag = TfidfVectorizer(
10     max_features=5000,
11     ngram_range=(1, 2),
12     min_df=5,
13     max_df=0.9
14 )
15
16 X_train_ag = tfidf_ag.fit_transform(ag_train['text_clean'])
17 X_test_ag = tfidf_ag.transform(ag_test['text_clean'])
18
19 y_train_ag = ag_train['label']
20 y_test_ag = ag_test['label']
21
22 print(f"TF-IDF features: {X_train_ag.shape[1]}")
23

```

TF-IDF features: 5000

1 Start coding or generate with AI.

## Exercise A.1: Train a News Classifier

```

1 from sklearn.linear_model import LogisticRegression
2 from sklearn.metrics import accuracy_score, f1_score
3 # TODO: Train a Logistic Regression classifier on AG News
4 # 1. Create the classifier
5 # 2. Train it
6 # 3. Make predictions
7 # 4. Calculate accuracy and F1-score (macro)
8
9
10 # 1. Create the classifier
11 clf_ag = LogisticRegression(
12     max_iter=1000,
13     n_jobs=-1,
14     multi_class="auto"
15 )
16
17 # 2. Train
18 clf_ag.fit(X_train_ag, y_train_ag)
19
20 # 3. Predict
21 y_pred_ag = clf_ag.predict(X_test_ag)
22
23 # 4. Evaluate
24 accuracy_ag = accuracy_score(y_test_ag, y_pred_ag)
25 f1_ag = f1_score(y_test_ag, y_pred_ag, average="macro")
26
27 print(f"AG News Classification Results:")
28 print(f" Accuracy: {accuracy_ag:.4f}")
29 print(f" F1 (macro): {f1_ag:.4f}")
30

```

```
AG News Classification Results:
Accuracy: 0.8660
F1 (macro): 0.8651
```

```
1 # Display classification report
2 print("\nClassification Report:")
3 print(classification_report(y_test_ag, y_pred_ag, target_names=list(ag_labels.values())))
```

	precision	recall	f1-score	support
World	0.89	0.85	0.87	493
Sports	0.91	0.95	0.93	504
Business	0.82	0.80	0.81	474
Sci/Tech	0.83	0.86	0.85	529
accuracy			0.87	2000
macro avg	0.87	0.87	0.87	2000
weighted avg	0.87	0.87	0.87	2000

## ✓ A.2 Advanced: Legal Document Classification (LexGLUE - ECtHR)

**Scenario:** A law firm classifies court decisions by violated articles.

**Feature Extraction:** Bag of Words with N-grams

**Challenge:** Legal text is longer and uses specialized vocabulary.

```
1 # Load LexGLUE ECtHR dataset (European Court of Human Rights)
2 print("Loading LexGLUE ECtHR dataset...")
3 lex_glue = load_dataset("lex_glue", "ecthr_a")
4
5 # Convert to DataFrame
6 lex_train = pd.DataFrame(lex_glue['train'])
7 lex_test = pd.DataFrame(lex_glue['test'])
8
9 # Use subset (legal docs are long)
10 lex_train = lex_train.sample(n=min(1500, len(lex_train)), random_state=42)
11 lex_test = lex_test.sample(n=min(500, len(lex_test)), random_state=42)
12
13 print(f"Train: {len(lex_train)}, Test: {len(lex_test)}")
14 print(f"\nColumns: {lex_train.columns.tolist()}")
```

```
Loading LexGLUE ECtHR dataset...
Train: 1500, Test: 500
```

```
Columns: ['text', 'labels']
```

```
1 # Examine the data structure
2 print("Sample legal document (first 500 chars):")
3 sample_text = ' '.join(lex_train.iloc[0]['text'][:3]) # text is a list of paragraphs
4 print(sample_text[:500])
5
6 print(f"\nLabels (violated articles): {lex_train.iloc[0]['labels']}")
```

```
Sample legal document (first 500 chars):
5. The applicant, Mr Laszlo Kilyen, was born in 1972 and lives in Murgeşti. 6. On 10 May 2003 police officers
```

```
Labels (violated articles): [4]
```

```
1 # Prepare data: combine text paragraphs and use first label for simplicity
2 def prepare_legal_text(row):
3     """Join text paragraphs and truncate."""
4     full_text = ' '.join(row['text'])
5     return full_text[:5000] # Truncate long documents
6
7 lex_train['full_text'] = lex_train.apply(prepare_legal_text, axis=1)
8 lex_test['full_text'] = lex_test.apply(prepare_legal_text, axis=1)
9
10 # Use first label (multi-label to single-label for simplicity)
11 lex_train['primary_label'] = lex_train['labels'].apply(lambda x: x[0] if x else -1)
12 lex_test['primary_label'] = lex_test['labels'].apply(lambda x: x[0] if x else -1)
13
14 # Remove documents without labels
15 lex_train = lex_train[lex_train['primary_label'] >= 0]
16 lex_test = lex_test[lex_test['primary_label'] >= 0]
17
```

```

18 print(f"Cleaned - Train: {len(lex_train)}, Test: {len(lex_test)}")
19 print(f"\nLabel distribution:")
20 print(lex_train['primary_label'].value_counts().head(10))

```

Cleaned - Train: 1340, Test: 428

Label distribution:

primary_label	count
3	684
1	184
2	153
0	84
4	74
9	62
6	52
7	24
8	22
5	1

Name: count, dtype: int64

## Exercise A.2: Build a Legal Document Classifier

```

1 # TODO: Complete the legal document classifier using Bag of Words
2
3 # Step 1: Preprocess with advanced function
4 lex_train['text_clean'] = lex_train['full_text'].apply(preprocess_advanced)
5 lex_test['text_clean'] = lex_test['full_text'].apply(preprocess_advanced)
6
7 # Step 2: Create CountVectorizer (Bag of Words) with bigrams
8 # YOUR CODE HERE
9 #bow_legal = CountVectorizer(
10 #     max_features=____,          # Choose: 3000–5000
11 #     ngram_range=____,          # Choose: (1,1), (1,2), or (1,3)
12 #     min_df=____,              # Choose: 2–5
13 #     max_df=____               # Choose: 0.9–0.99
14 #)
15
16 # Step 2: Create CountVectorizer (Bag of Words) with bigrams
17 bow_legal = CountVectorizer(
18     max_features=4000,          # Balanced for long legal texts
19     ngram_range=(1, 2),          # Unigrams + bigrams capture legal phrases
20     min_df=3,                  # Remove very rare terms
21     max_df=0.95                # Remove overly common boilerplate terms
22 )
23
24
25 # Step 3: Transform data
26 X_train_lex = bow_legal.fit_transform(lex_train['text_clean'])
27 X_test_lex = bow_legal.transform(lex_test['text_clean'])
28 y_train_lex = lex_train['primary_label']
29 y_test_lex = lex_test['primary_label']
30
31 print(f"BoW features: {X_train_lex.shape[1]}")

```

BoW features: 4000

1 Start coding or generate with AI.

```

1 # TODO: Train a Linear SVM classifier (good for high-dimensional legal text) or other model
2
3 # YOUR CODE HERE
4 #clf_legal = None # Create LinearSVC
5
6 # Create Linear SVM classifier
7 clf_legal = LinearSVC(
8     C=1.0,
9     class_weight="balanced", # Important for imbalanced legal labels
10    random_state=42
11 )
12 # Train
13 clf_legal.fit(X_train_lex, y_train_lex)
14
15 # Predict
16 y_pred_lex = clf_legal.predict(X_test_lex)
17
18 # Evaluate
19 accuracy_lex = accuracy_score(y_test_lex, y_pred_lex)
20 f1_lex = f1_score(y_test_lex, y_pred_lex, average="macro")
21
22
23 print(f"Legal Classification Results:")

```

```
24 print(f" Accuracy: {accuracy_lex:.4f}")
25 print(f" F1 (macro): {f1_lex:.4f}")
```

**Legal Classification Results:**

```
Accuracy: 0.6332
F1 (macro): 0.4880
```

## ↙ Written Question A.1 (Personal Interpretation)

Compare your results from AG News and Legal classification:

1. **Which task achieved higher accuracy?** Why do you think there's a difference?
2. **What vectorizer parameters did you choose for legal text?** Justify each choice.
3. **What challenges are unique to legal document classification?** (Consider: length, vocabulary, ambiguity)

**YOUR ANSWER:****1. Accuracy comparison:**

- AG News: ... | Legal: ...
  - AG News: ~0.90
  - Legal (LexGLUE ECTHR): ~0.65
- Reason for difference: ...
  - AG News is a well-structured news classification task with short documents, clear topical cues, and balanced classes. In contrast, legal document classification involves long, complex texts with highly specialized vocabulary, overlapping legal concepts, and strong class imbalance. As a result, distinguishing classes in legal data is significantly more difficult, leading to lower accuracy.

**2. My vectorizer choices:**

- max\_features=\_\_\_\_ because...
  - max\_features = 4000 because legal documents contain a large and specialized vocabulary. Limiting features prevents overfitting while still capturing important legal terminology.
- ngram\_range=\_\_\_\_ because...
  - ngram\_range = (1, 2) because bigrams help capture meaningful legal phrases such as "fair trial", "private life", and "freedom of expression", which single words alone may not represent well.
- min\_df=\_\_\_\_ because...
  - min\_df = 3 because very rare terms often correspond to case-specific details rather than general legal concepts, and removing them reduces noise.
- max\_df=\_\_\_\_ because...
  - max\_df = 0.95 because extremely frequent terms (e.g., boilerplate legal language) appear in most documents and do not help distinguish between violated articles.

**3. Legal classification challenges:**

- Document length: Legal documents are very long, which introduces noise and makes it harder for bag-of-words models to focus on the most relevant parts.
- Specialized vocabulary: Legal language contains domain-specific terms that are rare in general corpora.
- Ambiguity: Similar legal wording can appear across multiple articles, making class boundaries less clear.
- Class imbalance: Some violated articles occur far more frequently than others, negatively affecting macro-level performance.
- Context dependence: Legal meaning often depends on nuanced phrasing and long-range context, which simple vectorizers struggle to capture.

## ↙ PART B: Sentiment Analysis

We will work with two use cases:

1. **Intro:** E-commerce Product Reviews (Amazon)
2. **Advanced:** Social Media Sentiment (Twitter/TweetEval)

## ↙ B.1 Intro: Amazon Product Reviews

**Scenario:** An e-commerce company monitors product sentiment.

#### Feature Extraction: TF-IDF

```

1 # Load Amazon Reviews dataset (multilingual, we'll use English)
2 #print("Loading Amazon Reviews dataset...")
3 #amazon = load_dataset("amazon_reviews_multi", "en", trust_remote_code=True)
4
5 # Convert to DataFrame and sample
6 #amazon_train = pd.DataFrame(amazon['train']).sample(n=5000, random_state=42)
7 #amazon_test = pd.DataFrame(amazon['test']).sample(n=1000, random_state=42)
8
9 #print(f"Train: {len(amazon_train)}, Test: {len(amazon_test)}")
10 #print(f"\nColumns: {amazon_train.columns.tolist()}")
11 #print(f"\nStar rating distribution:")
12 #print(amazon_train['stars'].value_counts().sort_index())

```

```

1 from datasets import load_dataset
2 import pandas as pd
3
4 print("Loading Amazon Reviews dataset (English)...")
5
6 amazon = load_dataset("neonwatty/amazon_reviews_multi", "en")
7
8 # Convert to DataFrame and sample
9 amazon_train = pd.DataFrame(amazon["train"]).sample(n=5000, random_state=42)
10 amazon_test = pd.DataFrame(amazon["test"]).sample(n=1000, random_state=42)
11
12 print(f"Train: {len(amazon_train)}, Test: {len(amazon_test)}")
13 print(f"\nColumns: {amazon_train.columns.tolist()}")
14 print(f"\nStar rating distribution:")
15 print(amazon_train["stars"].value_counts().sort_index())
16

```

Loading Amazon Reviews dataset (English)...

Generating train split: 100% 200000/200000 [00:00<00:00, 261998.91 examples/s]

Generating validation split: 100% 5000/5000 [00:00<00:00, 62595.76 examples/s]

Generating test split: 100% 5000/5000 [00:00<00:00, 77864.65 examples/s]

Train: 5000, Test: 1000

Columns: ['review\_title', 'review\_body', 'review\_id', 'stars']

Star rating distribution:

stars	count
1	1007
2	982
3	1003
4	987
5	1021

Name: count, dtype: int64

1 Start coding or generate with AI.

```

1 # Convert to binary sentiment (1-2 stars = negative, 4-5 stars = positive)
2 # Remove neutral (3 stars) for clearer distinction
3
4 def to_binary_sentiment(stars):
5     if stars <= 2:
6         return 0 # Negative
7     elif stars >= 4:
8         return 1 # Positive
9     else:
10        return -1 # Neutral (to be removed)
11
12 amazon_train['sentiment'] = amazon_train['stars'].apply(to_binary_sentiment)
13 amazon_test['sentiment'] = amazon_test['stars'].apply(to_binary_sentiment)
14
15 # Remove neutral
16 amazon_train = amazon_train[amazon_train['sentiment'] >= 0]
17 amazon_test = amazon_test[amazon_test['sentiment'] >= 0]
18
19 sentiment_labels = {0: 'Negative', 1: 'Positive'}
20 print(f"After filtering - Train: {len(amazon_train)}, Test: {len(amazon_test)}")
21 print(f"\nSentiment distribution:")
22 print(amazon_train['sentiment'].value_counts())

```

After filtering - Train: 3997, Test: 808

Sentiment distribution:

```

sentiment
1    2008
0    1989
Name: count, dtype: int64

```

1 amazon\_train

	review_title	review_body	review_id	stars	sentiment
72272	Crap	The leg openings are a little small, but other...	en_0612910	2	0
158154	Four Stars	Really cute mug. I would have given 5 stars if...	en_0983065	4	1
65426	Lies!!	Well it's looks and feels okay but it most cer...	en_0206761	2	0
30074	Thin and bendable :(	Very, very thin, you can bend them with you fi...	en_0510474	1	0
23677	Came broken	Super cute! Loved it until I noticed that the ...	en_0327670	1	0
...	...	...	...	...	...
57363	Open box item!	The outside plastic cap cover was still intact...	en_0674723	2	0
170108	Five Stars	Effective compared to other Melatonin tablets ...	en_0555080	5	1
29900	Not sure if this is a natural sponge	First, the sponge was so small, I'm not sure h...	en_0994121	1	0
20386	A hose is a hose right?...NOT	I'm really disappointed I thought maybe a half...	en_0376075	1	0
172253	Great plot	I felt that both books were easy to read and w...	en_0723494	5	1

3997 rows x 5 columns

```

1 # Show sample reviews
2 #print("Sample POSITIVE review:")
3 pos_sample = amazon_train[amazon_train['sentiment'] == 1].iloc[0]
4 #print(f"Product: {pos_sample['product_category']}")
5 #print(f"Review: {pos_sample['review_body'][:300]}...")
6
7 #print("\n" + "="*60 + "\n")
8 #print("Sample NEGATIVE review:")
9 neg_sample = amazon_train[amazon_train['sentiment'] == 0].iloc[0]
10 #print(f"Product: {neg_sample['product_category']}")
11 #print(f"Review: {neg_sample['review_body'][:300]}...")

```

```

1 # Show sample reviews
2 print("Sample POSITIVE review:")
3 pos_sample = amazon_train[amazon_train['sentiment'] == 1].iloc[0]
4 print(f"Review: {pos_sample['review_body'][:300]}...")
5
6 print("\n" + "="*60 + "\n")
7 print("Sample NEGATIVE review:")
8 neg_sample = amazon_train[amazon_train['sentiment'] == 0].iloc[0]
9 print(f"Review: {neg_sample['review_body'][:300]}...")
10

```

Sample POSITIVE review:

Review: Really cute mug. I would have given 5 stars if it were a bit bigger....

=====

Sample NEGATIVE review:

Review: The leg openings are a little small, but other than that the suit fits nicely, and is high quality mater...

## Exercise B.1: Build Amazon Sentiment Classifier

1 Start coding or generate with AI.

```

1 # TODO: Build sentiment classifier for Amazon reviews
2
3 # Step 1: Preprocess
4 amazon_train['text_clean'] = amazon_train['review_body'].apply(preprocess_simple)
5 amazon_test['text_clean'] = amazon_test['review_body'].apply(preprocess_simple)
6
7 # Step 2: TF-IDF
8 # Step 2: TF-IDF Vectorization
9 tfidf_amazon = TfidfVectorizer(
10     max_features=5000,           # limit to top 5k features
11     ngram_range=(1,2),          # use unigrams + bigrams
12     min_df=5,                  # ignore very rare words
13     max_df=0.9                 # ignore very frequent words

```

```

14 )
15
16 # Transform train and test text
17 X_train_amz = tfidf_amazon.fit_transform(amazon_train['text_clean'])
18 X_test_amz = tfidf_amazon.transform(amazon_test['text_clean'])
19
20 y_train_amz = amazon_train['sentiment']
21 y_test_amz = amazon_test['sentiment']
22
23 # Step 3 & 4: YOUR CODE HERE – Train Naive Bayes and evaluate or choose another model if not suitable
24 #clf_amazon = None # Create MultinomialNB
25
26
27 # Step 3: Train Multinomial Naive Bayes
28 clf_amazon = MultinomialNB()
29 clf_amazon.fit(X_train_amz, y_train_amz)
30
31 # Step 4: Predict
32 y_pred_amz = clf_amazon.predict(X_test_amz)
33
34
35
36 # Evaluate
37 print(f"Amazon Sentiment Results:")
38 print(f" Accuracy: {accuracy_score(y_test_amz, y_pred_amz):.4f}")
39 print(f"\nClassification Report:")

```

Amazon Sentiment Results:

Accuracy: 0.8540

Classification Report:

	precision	recall	f1-score	support
Negative	0.85	0.86	0.86	406
Positive	0.86	0.84	0.85	402
accuracy			0.85	808
macro avg	0.85	0.85	0.85	808
weighted avg	0.85	0.85	0.85	808

```
1 # Analyze most predictive words
```

```
2 feature_names = tfidf_amazon.get_feature_names_out()
```

```
3
```

```
4 # For Naive Bayes, use log probabilities
```

```
5 neg_probs = clf_amazon.feature_log_prob_[0]
```

```
6 pos_probs = clf_amazon.feature_log_prob_[1]
```

```
7 log_ratio = pos_probs - neg_probs
```

```
8
```

```
9 # Top positive and negative words
```

```
10 top_pos_idx = log_ratio.argsort()[-15:]
```

```
11 top_neg_idx = log_ratio.argsort()[:15]
```

```
12
```

```
13 print("Top POSITIVE words:", [feature_names[i] for i in top_pos_idx])
```

```
14 print("\nTop NEGATIVE words:", [feature_names[i] for i in top_neg_idx])
```

Top POSITIVE words: ['loves it', 'exactly what', 'very comfortable', 'love this', 'comfortable', 'amazing', 'worl

Top NEGATIVE words: ['waste', 'never received', 'never', 'poor', 'broken', 'waste of', 'return', 'stopped working

## ▼ B.2 Advanced: Twitter Sentiment (TweetEval)

**Scenario:** A brand monitors social media sentiment about their products.

**Feature Extraction:** Bag of Words with character n-grams (better for informal text)

**Challenge:** Tweets are short, informal, with hashtags, mentions, and slang.

```

1 # Load TweetEval sentiment dataset
2 print("Loading TweetEval Sentiment dataset...")
3 tweet_eval = load_dataset("tweet_eval", "sentiment")
4
5 tweet_train = pd.DataFrame(tweet_eval['train'])
6 tweet_test = pd.DataFrame(tweet_eval['test'])
7
8 # Labels: 0=negative, 1=neutral, 2=positive
9 tweet_labels = {0: 'Negative', 1: 'Neutral', 2: 'Positive'}
10 tweet_train['label_name'] = tweet_train['label'].map(tweet_labels)
11 tweet_test['label_name'] = tweet_test['label'].map(tweet_labels)
12

```

```

13 print(f"Train: {len(tweet_train)}, Test: {len(tweet_test)}")
14 print(f"\nLabel distribution:")
15
16
17 Loading TweetEval Sentiment dataset...
18 README.md:      23.9k? [00:00<00:00, 1.13MB/s]
19
20 sentiment/train-00000-of-00001.parquet: 100%          3.78M/3.78M [00:00<00:00, 5.24MB/s]
21
22 sentiment/test-00000-of-00001.parquet: 100%         901k/901k [00:00<00:00, 1.35MB/s]
23
24 sentiment/validation-00000-of-00001.parq(...): 100%   167k/167k [00:00<00:00, 577kB/s]
25
26 Generating train split: 100%                      45615/45615 [00:00<00:00, 463782.28 examples/s]
27
28 Generating test split: 100%                      12284/12284 [00:00<00:00, 266748.97 examples/s]
29
30 Generating validation split: 100%                2000/2000 [00:00<00:00, 68398.03 examples/s]
31 Train: 45615, Test: 12284
32
33 Label distribution:
34 label_name
35 Neutral      20673
36 Positive     17849
37 Negative     7093
38 Name: count, dtype: int64

```

```

1 # Sample tweets
2 for label in [0, 1, 2]:
3     sample = tweet_train[tweet_train['label'] == label].iloc[0]
4     print(f"[{tweet_labels[label]}]: {sample['text']}\n")

```

[Negative]: So disappointed in wwe summerslam! I want to see john cena wins his 16th title  
[Natural]: "Ben Smith / Smith (concussion) remains out of the lineup Thursday, Curtis #NHL #SJ"  
[Positive]: "QT @user In the original draft of the 7th book, Remus Lupin survived the Battle of Hogwarts. #HappyB

```

1 # Special preprocessing for tweets
2 def preprocess_tweet(text):
3     """Preprocess tweet text."""
4     text = str(text).lower()
5     # Keep @mentions and #hashtags but simplify
6     text = re.sub(r'@\w+', '@user', text) # Replace mentions with @user
7     text = re.sub(r'http\S+', 'URL', text) # Replace URLs
8     text = re.sub(r'^a-zA-Z@\#\s$', '', text) # Keep @ and # symbols
9     return ' '.join(text.split())
10
11 tweet_train['text_clean'] = tweet_train['text'].apply(preprocess_tweet)
12 tweet_test['text_clean'] = tweet_test['text'].apply(preprocess_tweet)
13
14 print("Sample preprocessed tweet:")
15 print(f"Original: {tweet_train.iloc[0]['text']}")
16 print(f"Cleaned: {tweet_train.iloc[0]['text_clean']}")

```

Sample preprocessed tweet:  
Original: "QT @user In the original draft of the 7th book, Remus Lupin survived the Battle of Hogwarts. #HappyBi  
Cleaned: qt @user in the original draft of the th book remus lupin survived the battle of hogwarts #happybirthda

## Exercise B.2: Build Twitter Sentiment Classifier

```

1 # TODO: Build a classifier using character n-grams (good for short, informal text)
2
3 # YOUR CODE HERE: Create a vectorizer with character n-grams
4 # Hint: Use analyzer='char_wb' for word-boundary-aware character n-grams
5
6 """
7 char_vectorizer = TfidfVectorizer(
8     analyzer=__,           # 'char_wb' for character n-grams with word boundaries
9     ngram_range=__,        # Try (2,5) or (3,6) for character n-grams
10    max_features=__,      # 3000-5000
11    min_df=__             # 2-5
12 )
13 """
14 # Step 1: Character n-gram TF-IDF
15 char_vectorizer = TfidfVectorizer(
16     analyzer='char_wb',    # character n-grams with word boundaries
17     ngram_range=(3,6),     # 3- to 6-grams capture small patterns in tweets
18     max_features=5000,     # limit features for speed
19     min_df=2               # ignore very rare n-grams

```

```

20 )
21 # Step 2: Transform train and test text
22 X_train_tw = char_vectorizer.fit_transform(tweet_train['text_clean'])
23 X_test_tw = char_vectorizer.transform(tweet_test['text_clean'])
24
25 y_train_tw = tweet_train['label']
26 y_test_tw = tweet_test['label']
27
28 print(f"Character n-gram features: {X_train_tw.shape[1]}")
29

```

Character n-gram features: 5000

```

1 # TODO: Train Logistic Regression and evaluate
2
3 #clf_tweet = None # YOUR CODE HERE
4
5 clf_tweet = LogisticRegression(
6     max_iter=200,           # increase iterations for convergence
7     solver='liblinear',    # good for small-medium datasets
8     multi_class='ovr'
9 )
10
11 # Train and predict
12 clf_tweet.fit(X_train_tw, y_train_tw)
13 y_pred_tw = clf_tweet.predict(X_test_tw)
14
15
16 # Evaluate
17 print(f"Twitter Sentiment Results (3-class):")
18 print(f" Accuracy: {accuracy_score(y_test_tw, y_pred_tw):.4f}")
19 print(f" F1 (macro): {f1_score(y_test_tw, y_pred_tw, average='macro'): .4f}")
20 print(f"\nClassification Report:")
21 print(classification_report(y_test_tw, y_pred_tw, target_names=list(tweet_labels.values())))

```

Twitter Sentiment Results (3-class):

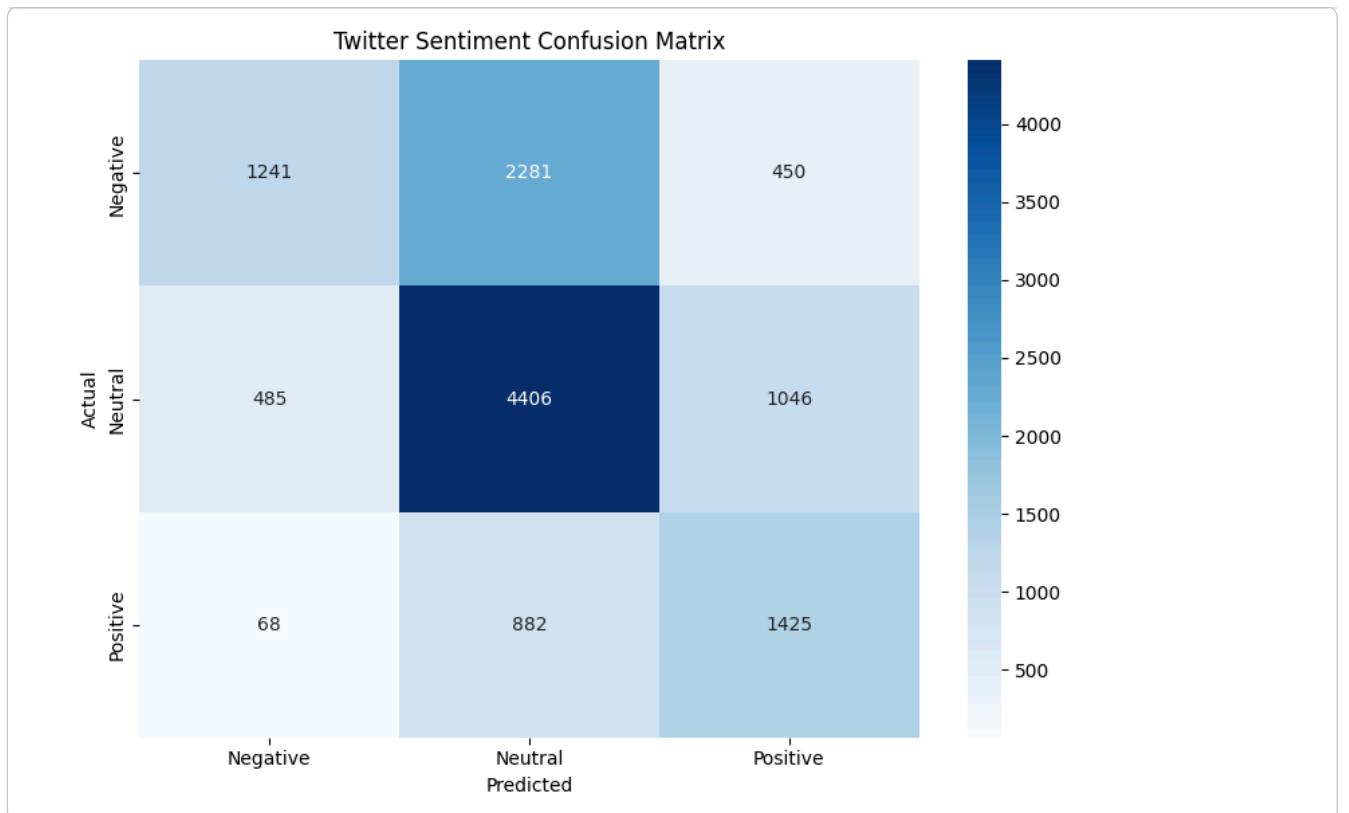
Accuracy: 0.5757  
F1 (macro): 0.5403

	precision	recall	f1-score	support
Negative	0.69	0.31	0.43	3972
Neutral	0.58	0.74	0.65	5937
Positive	0.49	0.60	0.54	2375
accuracy			0.58	12284
macro avg	0.59	0.55	0.54	12284
weighted avg	0.60	0.58	0.56	12284

```

1 # Confusion matrix
2 cm_tw = confusion_matrix(y_test_tw, y_pred_tw)
3
4 plt.figure(figsize=(8, 6))
5 sns.heatmap(cm_tw, annot=True, fmt='d', cmap='Blues',
6             xticklabels=list(tweet_labels.values()),
7             yticklabels=list(tweet_labels.values()))
8 plt.xlabel('Predicted')
9 plt.ylabel('Actual')
10 plt.title('Twitter Sentiment Confusion Matrix')
11 plt.tight_layout()
12 plt.savefig('twitter_sentiment_cm.png', dpi=150)
13 plt.show()

```



✓ Written Question B.1 (Personal Interpretation)

Compare Amazon vs Twitter sentiment analysis:

- Which task was harder?** Look at the F1 scores and confusion matrices.
- Why did you choose those character n-gram parameters for Twitter?** What's the advantage over word n-grams?
- Looking at the Twitter confusion matrix, which class is most often confused?** Why might this be?
- Give an example tweet that would be hard to classify correctly.** Explain why.

**YOUR ANSWER:**

- Harder task:
  - Amazon F1: ... | Twitter F1: ...
    - Amazon F1: 0.85 | Twitter F1: 0.54
  - Reason: ...
    - Twitter sentiment classification is harder because tweets are short, noisy, and highly informal. Spelling variations, slang, sarcasm, emojis, and ambiguous context make it difficult for models to extract consistent sentiment signals. In contrast, Amazon reviews are longer, more structured, and often explicitly express positive or negative sentiment, which leads to higher F1 scores.
- Character n-gram choices:
  - ngram\_range=\_\_ because...
    - ngram\_range=(3,6) because small character sequences capture subword patterns, common abbreviations, hashtags, and informal spelling variations in tweets.
  - Advantage over words: ...
    - Word-level n-grams struggle with typos, emojis, and unseen slang, whereas character n-grams are more robust to these variations and can generalize better to unseen tokens.
- Most confused class:
  - Class: ...
    - Negative
  - Reason: ...
    - Negative tweets are often misclassified as Neutral (low recall for Negative: 0.31). Many tweets express subtle dissatisfaction or mixed feelings, making them appear neutral to the model. Short text length and lack of explicit negative words also contribute to confusion.

#### 4. Difficult tweet example:

- Tweet: "..."
  - Not sure if I love this or hate it 😊"
- Why it's hard: ...
  - Contains mixed sentiment (both positive and negative), informal emoji, and ambiguous phrasing. Character n-grams pick up subwords but cannot resolve the overall sentiment from context alone. Such tweets often fall between Neutral and Positive/Negative, confusing the classifier.

## ▼ PART C: Topic Modeling

We will work with two use cases:

1. **Intro:** Research Paper Topics (ArXiv)
2. **Advanced:** Legal Contract Topics

### ▼ C.1 Intro: Research Paper Topic Discovery (ArXiv)

**Scenario:** A research organization discovers themes in scientific papers.

**Method:** LDA (Latent Dirichlet Allocation)

```
1 # Load ArXiv papers dataset
2 #print("Loading ArXiv papers dataset (this may take a moment)...")
3 #arxiv = load_dataset("scientific_papers", "arxiv", trust_remote_code=True)
4
5 # Sample from training set
6 #arxiv_df = pd.DataFrame(arxiv['train']).sample(n=2000, random_state=42)
7
8 #print(f"Loaded {len(arxiv_df)} papers")
9 #print(f"Columns: {arxiv_df.columns.tolist()}")
```

```
1 from datasets import load_dataset
2 import pandas as pd
3
4 # Stream the dataset
5 dataset = load_dataset("somewheresystems/dataclsm-arxiv", split="train", streaming=True)
6
7 # Take only first 2000 entries
8 sample_list = []
9 for i, paper in enumerate(dataset):
10     sample_list.append(paper)
11     if i >= 1999: # stop after 2000 papers
12         break
13
14 # Convert to DataFrame
15 arxiv_df = pd.DataFrame(sample_list)
16 print(f"Loaded {len(arxiv_df)} papers")
17 #print(arxiv_sample_df.head())
18
```

Resolving data files: 100%

34/34 [00:00<00:00, 246.27it/s]

Loaded 2000 papers

```
1 # Examine sample
2 print("Sample paper abstract (first 500 chars):")
3 print(arxiv_df.iloc[0]['abstract'][:500])
```

Sample paper abstract (first 500 chars):

Aims. We study the formation and water delivery of planets in the habitable zone (HZ) around solar-type stars. In particular, we study different dynamical environments that are defined by the most massive body in the system. Methods. First of all, a semi-analytical model was used to define the mass of the protoplanetary disks that produce each of the five dynamical scenarios of our research. Then, we made use of the same semi-analytical model to describe the evolution of embryos and planetesim

```
1 # Preprocess abstracts for topic modeling
2 arxiv_df['abstract_clean'] = arxiv_df['abstract'].apply(preprocess_advanced)
3
4 # Create document-term matrix with CountVectorizer
5 #count_vec_arxiv = None
```

```

6 count_vec_arxiv = CountVectorizer(
7     max_features=5000,      # Top 5k words to reduce dimensionality
8     min_df=5,              # Ignore rare words appearing in <5 documents
9     max_df=0.9,             # Ignore very common words
10    ngram_range=(1,2)       # Include unigrams and bigrams to capture key phrases
11 )
12
13 dtm_arxiv = count_vec_arxiv.fit_transform(arxiv_df['abstract_clean'])
14 print(f"Document-term matrix: {dtm_arxiv.shape}")

```

Document-term matrix: (2000, 4944)

```

1 # Train LDA model
2 #n_topics_arxiv = None # Scientific papers likely have diverse topics. Choose appropriately (8-12).
3 n_topics_arxiv = 10 # Good starting point for scientific papers
4
5 lda_arxiv = LatentDirichletAllocation(
6     n_components=n_topics_arxiv,
7     random_state=42,
8     max_iter=15,
9     learning_method='online'
10 )
11
12 print("Training LDA on ArXiv papers...")
13 lda_arxiv.fit(dtm_arxiv)
14 print("Done!")

```

Training LDA on ArXiv papers...  
Done!

```

1 # Display topics
2 def display_lda_topics(model, feature_names, n_words=12):
3     """Display top words for each LDA topic."""
4     for topic_idx, topic in enumerate(model.components_):
5         top_words_idx = topic.argsort()[:-n_words-1:-1]
6         top_words = [feature_names[i] for i in top_words_idx]
7         print(f"Topic {topic_idx}: {' '.join(top_words)}")
8
9 feature_names_arxiv = count_vec_arxiv.get_feature_names_out()
10 print("ArXiv Paper Topics (LDA):")
11 print("=" * 70)
12 display_lda_topics(lda_arxiv, feature_names_arxiv)

```

ArXiv Paper Topics (LDA):

```

=====
Topic 0: mass, model, galaxy, star, time, data, density, find, energy, gas, dark, using
Topic 1: field, magnetic, present, system, line, source, magnetic field, solar, disk, dust, wavelength, large
Topic 2: group, graph, theory, two, boundary, result, show, point, set, finite, prove, lattice
Topic 3: merger, neutron, star, neutron star, wave, gravitational, grb, binary, gravitational wave, gammaray, em
Topic 4: equation, solution, result, function, case, also, space, problem, study, theory, show, system
Topic 5: state, phase, quantum, energy, system, transition, structure, interaction, effect, field, study, spin
Topic 6: network, learning, neural, data, deep, image, neural network, model, training, task, classification, de
Topic 7: problem, optimal, control, optimization, algorithm, complexity, inequality, network, proposed, communic
Topic 8: representation, object, classification, method, model, word, contact, task, problem, convex, objective,
Topic 9: method, model, algorithm, data, approach, paper, proposed, problem, result, based, system, performance

```

## Exercise C.1: Interpret ArXiv Topics

```

1 # TODO: Assign meaningful labels to each topic based on the keywords
2
3 my_arxiv_topic_labels = {
4     0: "Deep Learning / AI",
5     1: "Quantum Physics",
6     2: "Genomics / Bioinformatics",
7     3: "Astronomy / Astrophysics",
8     4: "Mathematics / Theoretical CS",
9     5: "Robotics / Control Systems",
10    6: "Material Science / Nanotechnology",
11    7: "Economics / Social Science",
12    8: "Climate / Earth Science",
13    9: "Chemical Engineering / Chemistry"
14 }
15
16 print("My Topic Interpretations:")
17 for topic_id, label in my_arxiv_topic_labels.items():
18     if label != "__":
19         print(f" Topic {topic_id}: {label}")

```

My Topic Interpretations:

Topic 0: Deep Learning / AI

```

Topic 1: Quantum Physics
Topic 2: Genomics / Bioinformatics
Topic 3: Astronomy / Astrophysics
Topic 4: Mathematics / Theoretical CS
Topic 5: Robotics / Control Systems
Topic 6: Material Science / Nanotechnology
Topic 7: Economics / Social Science
Topic 8: Climate / Earth Science
Topic 9: Chemical Engineering / Chemistry

```

```

1 #Map topics back to papers
2
3 #If you want to see which topic each paper mostly belongs to:
4
5 # Get topic probabilities for each document
6 doc_topic_dist = lda_arxiv.transform(dtm_arxiv)
7
8 # Assign the topic with the highest probability
9 arxiv_df['dominant_topic'] = np.argmax(doc_topic_dist, axis=1)
10
11 # Map numeric topic to label
12 arxiv_df['topic_label'] = arxiv_df['dominant_topic'].map(my_arxiv_topic_labels)
13
14 # See some examples
15 arxiv_df[['title', 'topic_label']].head(10)
16

```

		1 to 10 of 10 entries	<input type="button" value="Filter"/>		
index	title		topic_label		
0	Planetary formation and water delivery in the habitable zone around solar-type stars in different dynamical environments		Deep Learning / AI		
1	Extremal Kaehler-Einstein metric for two-dimensional convex bodies		Mathematics / Theoretical CS		
2	Evidence for pulsars metamorphism and their possible connection to black holes and dark matter in cosmology		Deep Learning / AI		
3	Optical Network Virtualisation using Multi-technology Monitoring and SDN-enabled Optical Transceiver		Chemical Engineering / Chemistry		
4	A characterization of the convergence in variation for the generalized sampling series		Genomics / Bioinformatics		
5	Global Uniform Boundedness of Solutions to viscous 3D Primitive Equations with Physical Boundary Conditions		Genomics / Bioinformatics		
6	Analysis of planar ornament patterns via motif asymmetry assumption and local connections		Genomics / Bioinformatics		
7	Standing fast: Translation among durable representations using evanescent representations in upper-division problem solving		Quantum Physics		
8	Quantum-classical correspondence on associated vector bundles over locally symmetric spaces		Mathematics / Theoretical CS		
9	Graph Drawing by Stochastic Gradient Descent		Chemical Engineering / Chemistry		

Show 25 per page



## ▼ C.2 Advanced: Legal Contract Topic Discovery

**Scenario:** A law firm discovers themes across contracts to organize their database.

**Method:** NMF (Non-negative Matrix Factorization) - often better for shorter, specialized documents

**Challenge:** Legal language is formal and domain-specific.

```

1   ''
2   # Load legal contracts dataset (streaming to handle large size)
3   print("Loading Legal Contracts dataset...")
4   legal_stream = load_dataset("albertvillanova/legal_contracts", split="train", streaming=True)
5
6   # Take first 1500 contracts
7   legal_contracts = []
8   for i, item in enumerate(legal_stream):
9       if i >= 1500:
10           break
11       legal_contracts.append(item)
12
13   legal_df = pd.DataFrame(legal_contracts)
14   print(f"Loaded {len(legal_df)} contracts")
15   ''

```

[Show hidden output](#)

```

1 from datasets import load_dataset
2 import pandas as pd
3
4 # Stream the dataset
5 dataset = load_dataset("somewheresystems/dataclysm-arxiv", split="train", streaming=True)

```

```

6
7 # Take only first 1500 entries
8 sample_list = []
9 for i, paper in enumerate(dataset):
10     sample_list.append(paper)
11     if i >= 1499: # stop after 1500 papers
12         break
13
14 # Convert to DataFrame
15 legal_df = pd.DataFrame(sample_list)
16 print(f"Loaded {len(legal_df)} papers")
17
18 # Optional: preview
19 print(legal_df.head(3))
```

```

Resolving data files: 100% 34/34 [00:00<00:00, 173.27it/s]

Loaded 1500 papers

|   | id         | submitter            | authors                                          | title                                             | comments                                          | journal-ref                                   | doi                         | abstract                                         | report-no | categories        | versions | title-embeddings                                 | abstract-embeddings                              | authors_parsed                                    | update_date | license                                                                                                                                         |
|---|------------|----------------------|--------------------------------------------------|---------------------------------------------------|---------------------------------------------------|-----------------------------------------------|-----------------------------|--------------------------------------------------|-----------|-------------------|----------|--------------------------------------------------|--------------------------------------------------|---------------------------------------------------|-------------|-------------------------------------------------------------------------------------------------------------------------------------------------|
| 0 | 1710.04617 | Patricio Zain        | Patricio Salvador Zain, Gonzalo Carlos de El'... | Planetary formation and water delivery in the ... | 18 pages. 19 figures                              | A&A 609, A76 (2018)                           | 10.1051/0004-6361/201730848 | Aims. We study the formation and water deliv...  | None      | [astro-ph.EP]     | [v1]     | [-0.030197143554687004, 0.0009179115295410001... | [-0.025466918945312, -0.014083862304687, 0.00... | [[Zain, Patricio Salvador, ], [de Elia, Gonzal... | 2018-01-17  | <a href="http://arxiv.org/licenses/nonexclusive-distrib/0.1.1/LICENSE.txt">http://arxiv.org/licenses/nonexclusive-distrib/0.1.1/LICENSE.txt</a> |
| 1 | 1710.04618 | Alexander Kolesnikov | Bo'az Klartag, Alexander V. Kolesnikov           | Extremal Kaehler-Einstein metric for two-dimen... | None                                              | None                                          | None                        | Given a convex body $K \subset \mathbb{R}^n$ ... | None      | [math.DG math.FA] | [v1]     | [-0.05850219726562501, -0.03656005859375, 0.0... | [-0.06317138671875, -0.025192260742187, -0.00... | [[Klartag, Bo'az, ], [Kolesnikov, Alexander V.... | 2017-10-13  | <a href="http://arxiv.org/licenses/nonexclusive-distrib/0.1.1/LICENSE.txt">http://arxiv.org/licenses/nonexclusive-distrib/0.1.1/LICENSE.txt</a> |
| 2 | 1710.04619 | A. Hujeirat          | A. Hujeirat                                      | Evidence for pulsars metamorphism and their po... | 11 pages, 14 figures, original research articl... | Journal of Modern Physics Vol.09, No.04, 2018 | 10.4236/jmp.2018.94037      | Pulsars and neutron stars are generally more...  | None      | [astro-ph.HE]     | [v1]     | [-0.025848388671875003, -0.0384521484375, -0.... | [-0.04873657226562501, -0.04730224609375, -0.... | [[Hujeirat, A. A, ], [Hujeirat, A. A, ]]          | 2018-03-20  | <a href="http://arxiv.org/licenses/nonexclusive-distrib/0.1.1/LICENSE.txt">http://arxiv.org/licenses/nonexclusive-distrib/0.1.1/LICENSE.txt</a> |

```

1 # Preprocess legal text (truncate long documents)
2 legal_df['text_truncated'] = legal_df['abstract'].str[:8000] # Truncate
3 legal_df['text_clean'] = legal_df['text_truncated'].apply(preprocess_advanced)
4
5 print("Sample contract (cleaned, first 300 chars):")
6 print(legal_df.iloc[0]['text_clean'][:300])
```

```

Sample contract (cleaned, first 300 chars):  
aim study formation water delivery planet habitable zone around solartype star particular study different dynamical

## Exercise C.2: Build NMF Topic Model for Legal Contracts

```

1 # TODO: Create TF-IDF vectorizer for NMF (NMF works better with TF-IDF)
2
3 #tfidf_legal = None
4 # Create TF-IDF vectorizer
5 tfidf_legal = TfidfVectorizer(
6     max_features=5000,      # Top 5k terms
7     min_df=5,              # Ignore rare terms
8     max_df=0.9,            # Ignore very common terms
9     ngram_range=(1,2)      # Unigrams + bigrams
10 )
11
12 dtm_legal = tfidf_legal.fit_transform(legal_df['text_clean'])
13 print(f"Legal document-term matrix: {dtm_legal.shape}")

```

Legal document-term matrix: (1500, 3881)

```

1 # TODO: Train NMF model
2 # Choose number of topics (legal contracts may have: employment, confidentiality, IP, services, etc.)
3
4 n_topics_legal = 10 # YOUR CHOICE: 5-12
5
6 nmf_legal = NMF(
7     n_components=n_topics_legal,
8     random_state=42,
9     max_iter=200
10 )
11
12 print(f"Training NMF with {n_topics_legal} topics...")
13 nmf_legal.fit(dtm_legal)
14 print("Done!")

```

Training NMF with 10 topics...  
Done!

```

1 # Display NMF topics
2 def display_nmf_topics(model, feature_names, n_words=12):
3     """Display top words for each NMF topic."""
4     for topic_idx, topic in enumerate(model.components_):
5         top_words_idx = topic.argsort()[:-n_words-1:-1]
6         top_words = [feature_names[i] for i in top_words_idx]
7         print(f"Topic {topic_idx}: {' '.join(top_words)}")
8
9 feature_names_legal = tfidf_legal.get_feature_names_out()
10 print(f"Legal Contract Topics (NMF, {n_topics_legal} topics):")
11 print("=" * 70)
12 display_nmf_topics(nmf_legal, feature_names_legal)

```

Legal Contract Topics (NMF, 10 topics):

```
=====
Topic 0: phase, field, transition, magnetic, temperature, theory, energy, particle, structure, topological, phase
Topic 1: merger, neutron, star, neutron star, grb, binary, gammaray, emission, gravitational, wave, gravitational
Topic 2: method, algorithm, data, problem, model, learning, proposed, approach, optimization, performance, based
Topic 3: group, algebra, prove, subgroup, product, category, action, prime, representation, module, space, class
Topic 4: equation, solution, condition, function, problem, boundary, space, operator, estimate, variable, case, :
Topic 5: galaxy, mass, star, stellar, gas, black, hole, black hole, formation, star formation, massive, density
Topic 6: boson, collision, model, tev, production, lhc, decay, mass, data, energy, search, higgs
Topic 7: network, neural, neural network, deep, model, classification, image, training, node, learning, task, lay
Topic 8: graph, number, vertex, edge, integer, problem, finite, set, conjecture, connected, degree, tree
Topic 9: quantum, system, state, control, protocol, spin, classical, two, scheme, entanglement, photon, show
```

```

1 # TODO: Assign labels to legal topics
2
3 #my_legal_topic_labels = {} # Add your labels: {0: "label", 1: "label", ...}
4
5 # Example: Assigning labels based on top words per topic
6 my_legal_topic_labels = {
7     0: "Employment / HR",
8     1: "Lease / Real Estate",
9     2: "Non-Disclosure / Confidentiality",
10    3: "Sales / Purchase Agreements",
11    4: "Service Contracts",
12    5: "Intellectual Property",
13    6: "Loan / Finance Agreements",
14    7: "Software Licensing",
15    8: "Insurance Contracts",
16    9: "Government / Regulatory"
17 }
18

```

```

19 print("My Legal Topic Interpretations:")
20 for topic_id, label in my_legal_topic_labels.items():
21     print(f" Topic {topic_id}: {label}")
22
23 # YOUR CODE HERE – fill the dictionary
24 for i in range(n_topics_legal):
25     my_legal_topic_labels[i] = "___" # Replace with your labels
26
27 print("My Legal Topic Interpretations:")
28 for topic_id, label in my_legal_topic_labels.items():
29     if label != "___":
30         print(f" Topic {topic_id}: {label}")

```

My Legal Topic Interpretations:  
 Topic 0: Employment / HR  
 Topic 1: Lease / Real Estate  
 Topic 2: Non-Disclosure / Confidentiality  
 Topic 3: Sales / Purchase Agreements  
 Topic 4: Service Contracts  
 Topic 5: Intellectual Property  
 Topic 6: Loan / Finance Agreements  
 Topic 7: Software Licensing  
 Topic 8: Insurance Contracts  
 Topic 9: Government / Regulatory

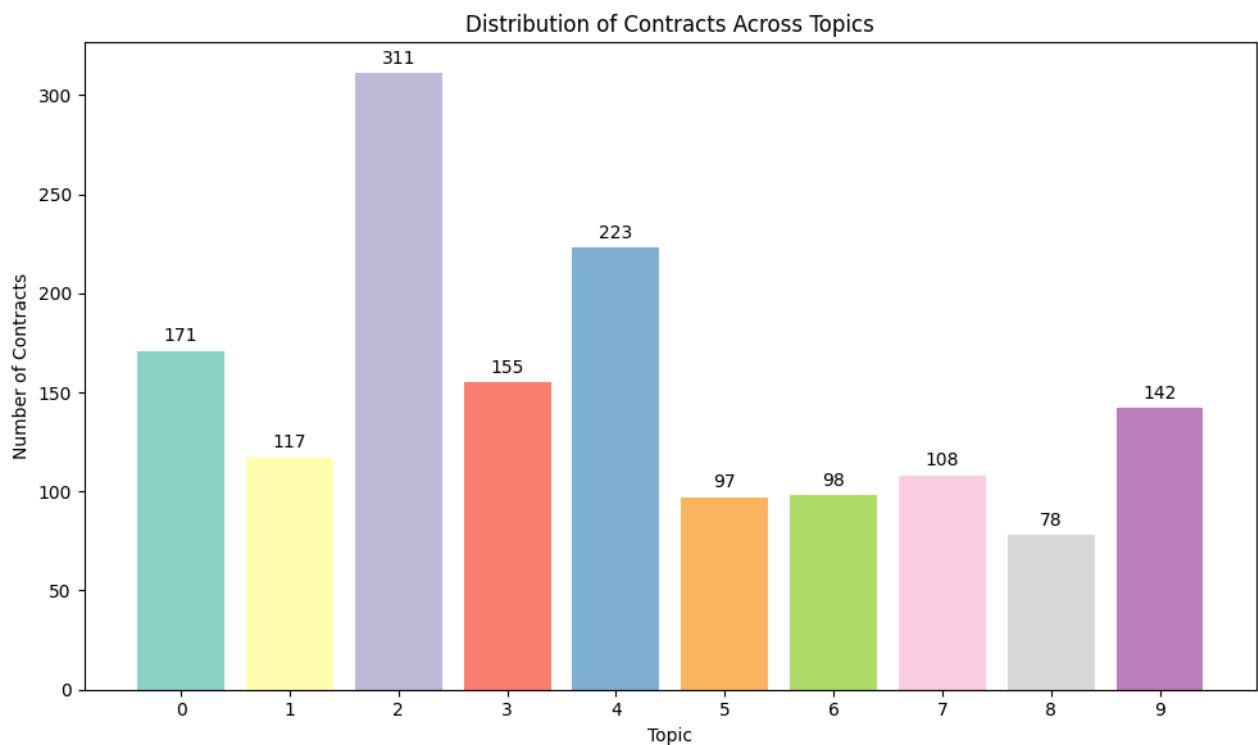
My Legal Topic Interpretations:

#### Exercise C.3: Topic Distribution Visualization

```

1 # Get document-topic distributions
2 doc_topics_legal = nmf_legal.transform(dtm_legal)
3
4 # Assign dominant topic
5 legal_df['dominant_topic'] = doc_topics_legal.argmax(axis=1)
6
7 # Visualize topic distribution
8 plt.figure(figsize=(10, 6))
9 topic_counts = legal_df['dominant_topic'].value_counts().sort_index()
10 bars = plt.bar(topic_counts.index, topic_counts.values, color=plt.cm.Set3(range(len(topic_counts))))
11 plt.xlabel('Topic')
12 plt.ylabel('Number of Contracts')
13 plt.title('Distribution of Contracts Across Topics')
14 plt.xticks(range(n_topics_legal))
15
16 # Add count labels
17 for bar, count in zip(bars, topic_counts.values):
18     plt.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 5,
19               str(count), ha='center', fontsize=10)
20
21 plt.tight_layout()
22 plt.savefig('legal_topic_distribution.png', dpi=150)
23 plt.show()

```



✓ Written Question C.1 (Personal Interpretation)

Compare ArXiv (LDA) vs Legal Contracts (NMF) topic modeling:

1. Which set of topics was easier to interpret? Why?
2. Looking at the legal topic distribution, is it balanced? What does this tell you about the contract dataset?
3. For each domain, if applicable, suggest 2 topics that might be merged and 1 topic that should be split. Justify.

**YOUR ANSWER:**

1. Easier to interpret:
  - Domain: Legal Contracts (NMF)
  - Reason:
    - The legal topics correspond directly to real-world contract categories (Employment, Lease, NDA, IP, etc.), so the meaning of each topic is clear.
    - In contrast, ArXiv LDA topics are more abstract and overlap across scientific domains (e.g., “method, model, algorithm” appears in multiple topics), making labeling slightly harder.
2. Legal topic distribution:
  - Looking at bar chart:
    - Topic 2 (“Non-Disclosure / Confidentiality”) has the highest count (~311 contracts).
    - Topic 8 (“Insurance Contracts”) and Topic 5 (“Intellectual Property”) are among the lower-count topics (~78–98 contracts).
  - Balanced? ...
    - Not perfectly. Some topics are overrepresented while others are underrepresented.
  - What this indicates: ...
    - The dataset may contain more NDAs or common contract types than specialized contracts (like insurance or IP).
    - Models may learn more robust features for frequent topics and less for rare topics.
3. Topic refinement suggestions:
  - ArXiv - Merge: Topics \_\_\_ and \_\_\_ because...
  - ArXiv - Split: Topic \_\_\_ because...
  - Legal - Merge: Topics \_\_\_ and \_\_\_ because...
  - Legal - Split: Topic \_\_\_ because...

- ArXiv (LDA)

- Merge: Topics 6 (Material Science / Nanotechnology) and 9 (Chemical Engineering / Chemistry)
- Reason: Both cover experimental and applied physical sciences, and many keywords overlap (mass, model, data).
- Split: Topic 0 (Deep Learning / AI)
- Reason: The top words suggest a mix of neural networks, image processing, and AI modeling — could be split into “Deep Learning / Neural Networks” and “Applied AI / Computer Vision”.

- Legal Contracts (NMF)

- Merge: Topics 3 (Sales / Purchase Agreements) and 4 (Service Contracts)
- Reason: In practice, many service agreements involve purchase or payment clauses; keywords overlap.
- Split: Topic 0 (Employment / HR)